

CS 199 BD Homework 3

Richard Lee, Raj Ramamurthy, Jay Bensal
NetIDs: rlee46, rmmrthy2 , bensal2

April 8, 2014

1 Linear Regression

Upon building a linear regression, we noticed that there are a lot of data points which were missing values. In order to build our linear regression, we removed these values.

1.1 R Code - Linear Regression

```
#read the CSV file
crime<-read.csv('communities.data', header=FALSE)

#remove nonpredictive variables
crime<-crime[c(-1,-2,-3,-4,-5)]
crime<-crime[sample(nrow(crime)),]

#replace all '?' with NA
crime[crime == '?'] <- NA

#only take the variables w/ all the values
drop_cols <- crime[complete.cases(crime), ]

#load Data Analysis and Graphics library
library(DAAG)

#perform multivariate linear regression
fit<-lm(V128 ~ V6+ V7+ V8+ V9+ V10+ V11+ V12+ V13+ V14+
  V15+ V16+ V17+ V18+ V19+ V20+ V21+ V22+ V23+ V24+ V25+
  V26+ V27+ V28+ V29+ V30+ V32+ V33+ V34+ V35+ V36+ V37
  + V38+ V39+ V40+ V41+ V42+ V43+ V44+ V45+ V46+ V47+
  V48+ V49+ V50+ V51+ V52+ V53+ V54+ V55+ V56+ V57+ V58+
  V59+ V60+ V61+ V62+ V63+ V64+ V65+ V66+ V67+ V68+ V69
  + V70+ V71+ V72+ V73+ V74+ V75+ V76+ V77+ V78+ V79+
  V80+ V81+ V82+ V83+ V84+ V85+ V86+ V87+ V88+ V89+ V90+)
```

```
V91+ V92+ V93+ V94+ V95+ V96+ V97+ V98+ V99+ V100+
V101+ V119+ V120+ V121+ V126, data=drop_cols)
```

```
#perform k-fold cross validation
cv.lm(df=drop_cols, fit, m=4)
```

1.2 Linear Regression Plot

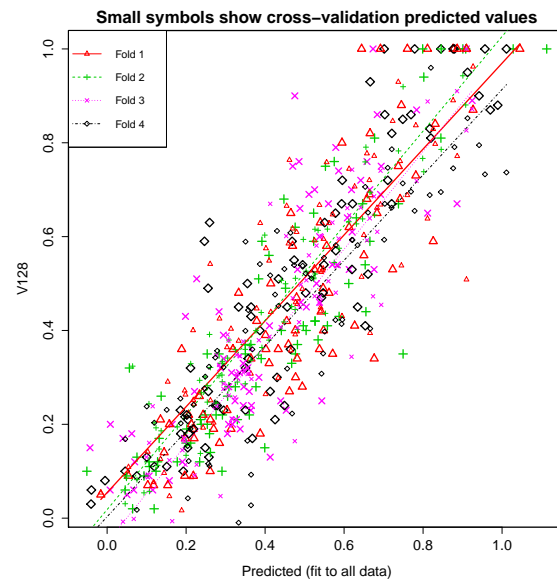


Figure 1: This graph shows the predicted fit across four different cross-validation folds of our data. Points with a good fit are close to the line $x = y$. Points that are further away from the line were predicted poorly.

Here is a plot that we made.

1.3 BoxCox Transformation - R code

```
#load in MASS library
library('MASS')
#pick a lambda value based on BoxCox
boxcox(fit, lambda = seq(0, 1, 1/10), plotit=TRUE)
```

1.4 Boxcox plot




Figure 2: We picked a value of $\lambda = 0.3$ based on this boxcox graph.

1.5 Applying the BoxCox Transformation

```
crime_lambda1 = drop_cols

#apply the boxcox transformation to the data
for (i in 1:nrow(crime_lambda1)){
  crime_lambda1[i, ncol(crime_lambda1)] = (crime_lambda1[i, ncol(crime_lambda1)]^0.3 - 1)/0.3
}

#perform multivariate linear regression based on new data
fit2 <- lm(V128 ~ V6+ V7+ V8+ V9+ V10+ V11+ V12+ V13+ V14+
  V15+ V16+ V17+ V18+ V19+ V20+ V21+ V22+ V23+ V24+ V25+
  V26+ V27+ V28+ V29+ V30+ V32+ V33+ V34+ V35+ V36+ V37
  + V38+ V39+ V40+ V41+ V42+ V43+ V44+ V45+ V46+ V47+
  V48+ V49+ V50+ V51+ V52+ V53+ V54+ V55+ V56+ V57+ V58+
  V59+ V60+ V61+ V62+ V63+ V64+ V65+ V66+ V67+ V68+ V69
  + V70+ V71+ V72+ V73+ V74+ V75+ V76+ V77+ V78+ V79+
  V80+ V81+ V82+ V83+ V84+ V85+ V86+ V87+ V88+ V89+ V90+
  V91+ V92+ V93+ V94+ V95+ V96+ V97+ V98+ V99+ V100+
  V101+ V119+ V120+ V121+ V126, data=crime_lambda1)

#perform k-fold cross validation on the new model
cv.lm(df=crime_lambda1, fit2, m=4)
```

1.6 BoxCox Graph

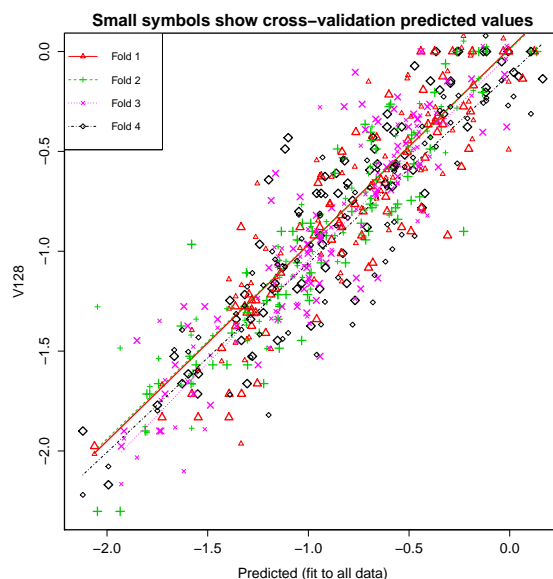


Figure 3: This graph shows the predicted fit across four different cross-validation folds of our data after performing the boxcox transformation. Points with a good fit are close to the line $x = y$. Points that are further away from the line were predicted poorly.

2 KNN

```
comm<-read.csv('communities.data', header=FALSE);  
  
#delete the first 5 vars which are not predictive  
comm<-comm[c(-1,-2,-3,-4,-5)]  
  
#load FNN library  
library('FNN')  
  
#replace all '?' with NA  
comm[comm == '?'] <- NA  
  
#then only take the ones w/ all the values  
full <- comm[complete.cases(comm),]  
lapply(full, as.numeric)
```

```

full = subset(full, select=c(V31, V102, V103, V104, V105,
, V106, V107, V108, V109, V111, V110, V112, V113, V114
, V115, V116, V117, V118, V122, V123, V124, V125, V127
))
comm_full = subset(comm, select=c(V31, V102, V103, V104,
V105, V106, V107, V108, V109, V111, V110, V112, V113,
V114, V115, V116, V117, V118, V122, V123, V124, V125,
V127))

#Perform nearest neighbor search with 21 neighbors
wtrain <- full[1:100, 1:(ncol(full)-1)]
wtrl <- full[1:100, (ncol(full))]
wtest <- full[101:200, 1:(ncol(full)-1)]
wtel <- full[101:200, ncol(full)]
results = knn(wtrain, wtest, wtrl, k = 21, algorithm="
cover_tree")

plot(as.numeric(results), as.numeric(wtel))

```

2.1 KNN Graph




Figure 4: Plot two

3 Impute

```

#load impute library
library(impute)

```

```

#impute the data
imputed=impute.knn(as.matrix(comm_full), k=10)

#perform multivariate linear regression
fit<-lm(V128 ~ V6+ V7+ V8+ V9+ V10+ V11+ V12+ V13+ V14+
  V15+ V16+ V17+ V18+ V19+ V20+ V21+ V22+ V23+ V24+ V25+
  V26+ V27+ V28+ V29+ V30+ V32+ V33+ V34+ V35+ V36+ V37
  + V38+ V39+ V40+ V41+ V42+ V43+ V44+ V45+ V46+ V47+
  V48+ V49+ V50+ V51+ V52+ V53+ V54+ V55+ V56+ V57+ V58+
  V59+ V60+ V61+ V62+ V63+ V64+ V65+ V66+ V67+ V68+ V69
  + V70+ V71+ V72+ V73+ V74+ V75+ V76+ V77+ V78+ V79+
  V80+ V81+ V82+ V83+ V84+ V85+ V86+ V87+ V88+ V89+ V90+
  V91+ V92+ V93+ V94+ V95+ V96+ V97+ V98+ V99+ V100+
  V101+ V119+ V120+ V121+ V126, data=as.data.frame(
  imputed$data))

#perform k-fold cross validation
cv.lm(df=as.data.frame(imputed$data), fit, m=4)

```

3.1 Impute plot




Figure 5: This graph shows the predicted fit across four different cross-validation folds of our imputed data. Points with a good fit are close to the line $x = y$. Points that are further away from the line were predicted poorly.

4 Modified Nearest Neighbor Search